

1/30 12:58:29 \*\*\*



# Spatial DBs

# Objectives/TOC

- spatial DBs: definition, characteristics, need, creation..
- spatial datatypes
- spatial operators
- spatial indices
- implementations
- miscellany

# What is a spatial database?

"A spatial database is a database that is optimized to store and query data related to objects in space, including points, lines and polygons."

In other words, it includes objects that have a SPATIAL location (and extent). A chief category of spatial data is geospatial data – derived from the geography of our earth.

Characteristics of geographic data:

- has location
- has size
- is auto-correlated
- scale dependent
- might be temporally dependent too

Geographic data is NOT 'business as usual'!

---

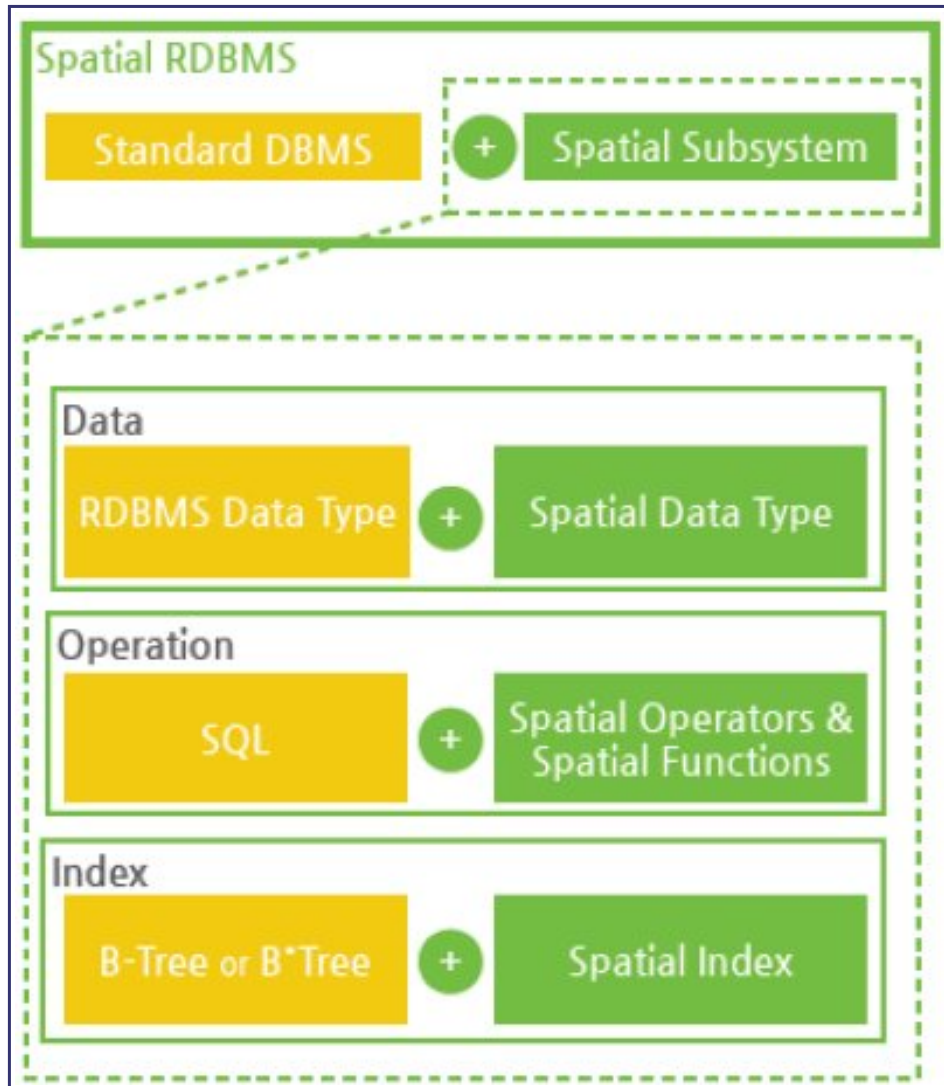
## Entity view vs field view

In spatial data analysis, we distinguish between two conceptions of space:

- entity view: space as an area filled with a set of discrete objects
- field view: space as an area covered with essentially continuous surfaces

For our purposes, we will adopt the 'entity' view, where space is populated by discrete objects (roads, buildings, rivers..).

# Components




So a spatial DB is a collection of the following, specifically built to handle spatial data:

- types
- operators
- indices

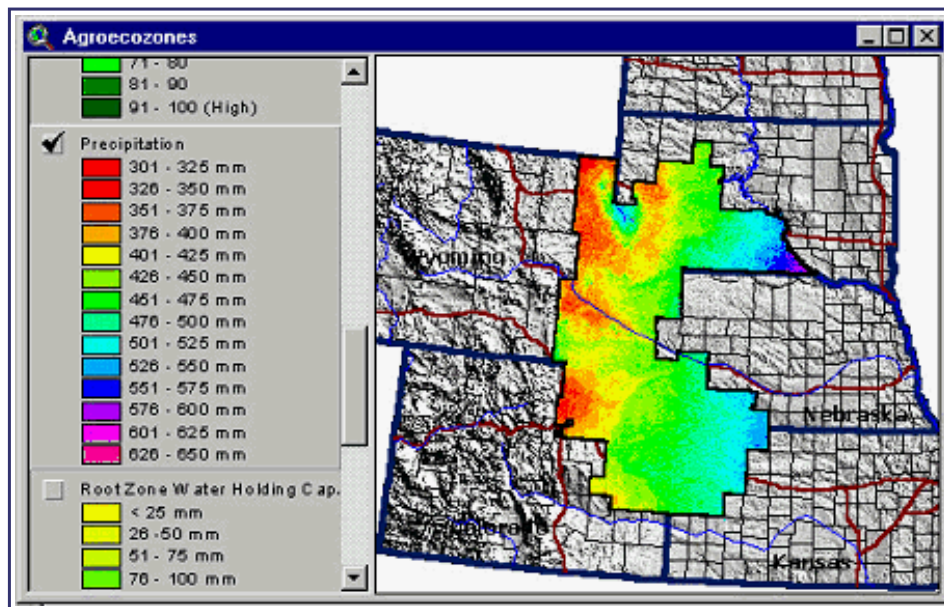
Soon, we will explore what types, operators and indices mean.

# Examples of spatial data

CAD data:

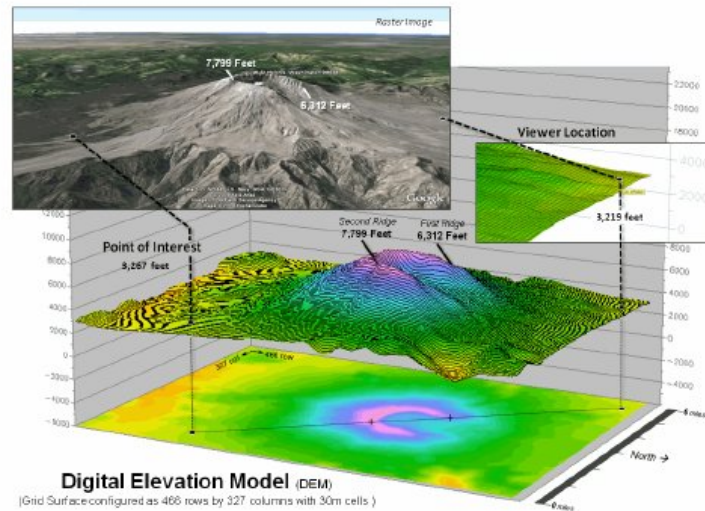


Agricultural data:



## 3D data:

- Three-dimensional data examples
  - Weather
  - Cartesian coordinates (3-D)
  - Topological
  - Satellite images



# What can be plotted on to a map?

- crime data
- spread of disease, risk of disease [look at this too]
- drug overdoses – over time
- census data
- income distribution, home prices
- locations of Starbucks (!)
- (real-time) traffic
- agricultural land use, deforestation

## Who creates/uses spatial data?

- **Army Field Commander:** Has there been any significant enemy troop movement since last night?
- **Insurance Risk Manager:** Which homes are most likely to be affected in the next great flood on the Mississippi?
- **Medical Doctor:** Based on this patient's MRI, have we treated somebody with a similar condition ?
- **Molecular Biologist:** Is the topology of the amino acid biosynthesis gene in the genome found in any other sequence feature map in the database ?
- **Astronomer:** Find all blue galaxies within 2 arcmin of quasars.

Various government agencies routinely coordinate spatial data collection and use, operating in effect, a national spatial data infrastructure (NSDI) – these include federal, state and local agencies. At the federal level, participating agencies include:

- Department of Commerce
  - Bureau of the Census
  - NIST
  - NOAA
- Department of Defense
  - Army Corps of Engineers
  - Defense Mapping Agency
- Department of the Interior
  - Bureau of Land Management



- Fish and Wildlife Service
  - U.S Geological Survey [earthquakes, map projections]
- Department of Agriculture
  - Agricultural Stabilization and Conservation Service
  - Economic Research Service
  - Forest Service
  - National Agriculture Statistical Service
  - Soil Conservation Service
- Department of Transportation
  - Federal Highway Administration
- Environmental Protection Agency
- NASA

As you can see, spatial data is a SERIOUS resource, vital to US' national interests.

# Where does spatial data come from?


Spatial data is created in a variety of ways:


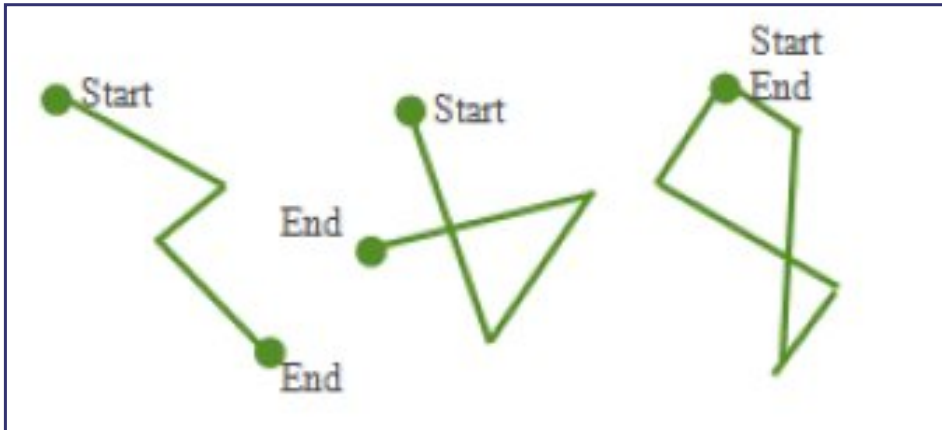
- CAD: user creation
- CAD: reverse engineering
- maps: cartography (surveying, plotting)
- maps: satellite imagery
- maps: 'copter, drone imagery
- maps: driving around
- maps: walking around

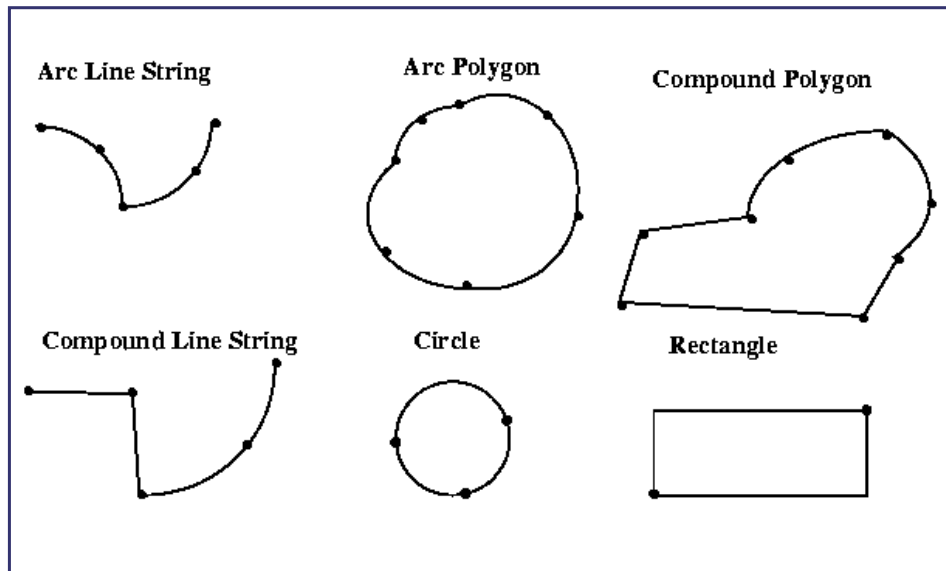
# What to store?

All spatial data can be described via the following entities/types:


- points/vertices/nodes
- polylines/arcs/linestrings
- polygons/regions
- pixels/raster







## Topological Elements and Relationships

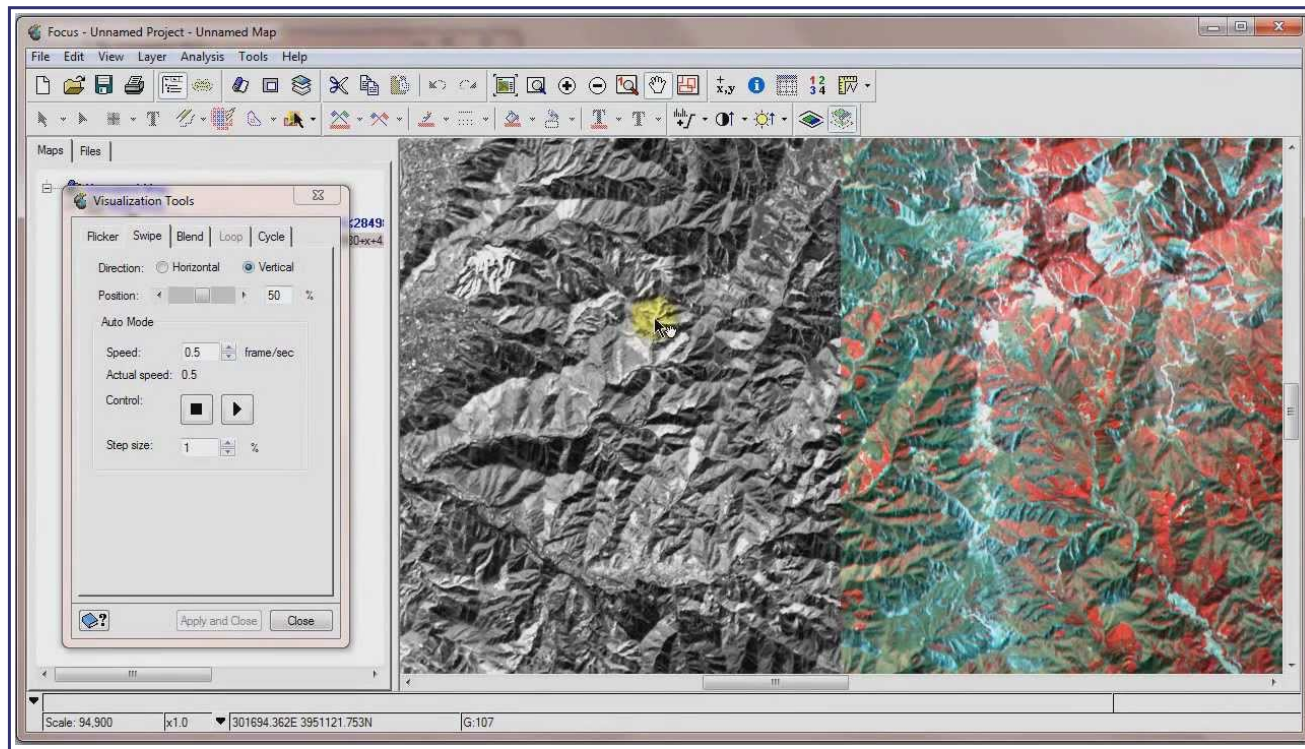


A Face

1 Edge

3 Node


→ Direction of edge



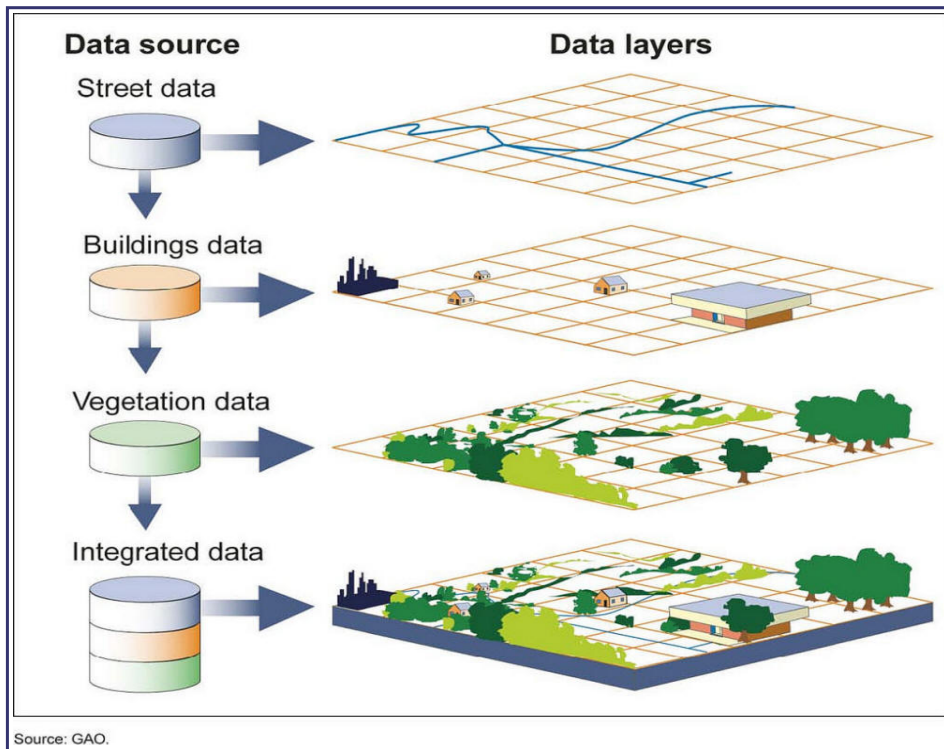
# Points, lines, polys => models and non-spatial attrs

Once we have spatial data (points, lines, polygons), we can:

- 'model' features such as lakes, soil type, highways, buildings etc, using the geometric primitives as underlying types
- add 'extra', non-spatial attributes/features to the underlying spatial data

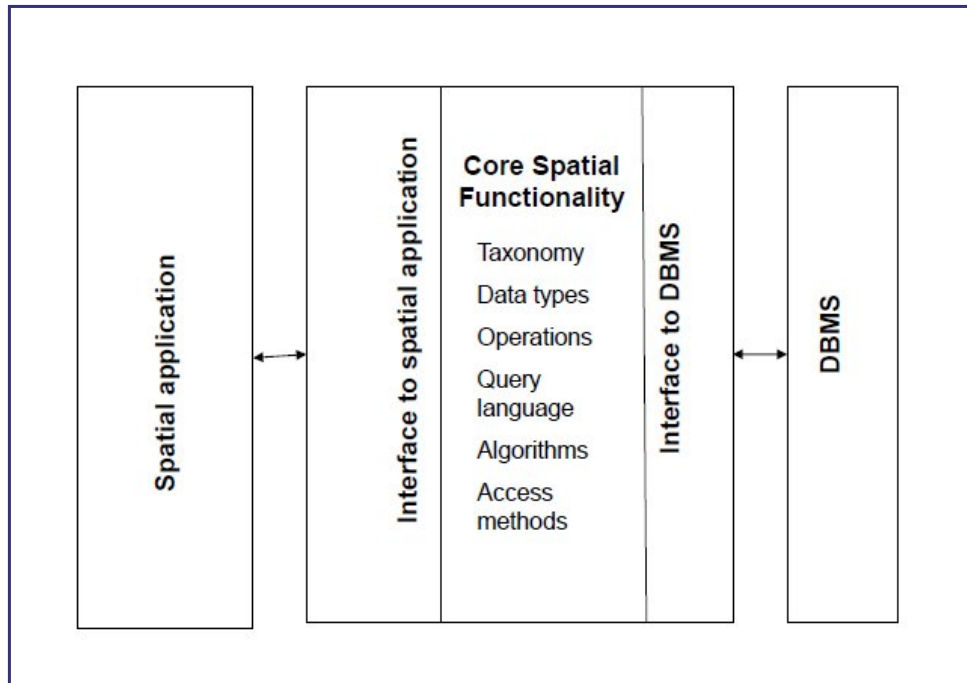


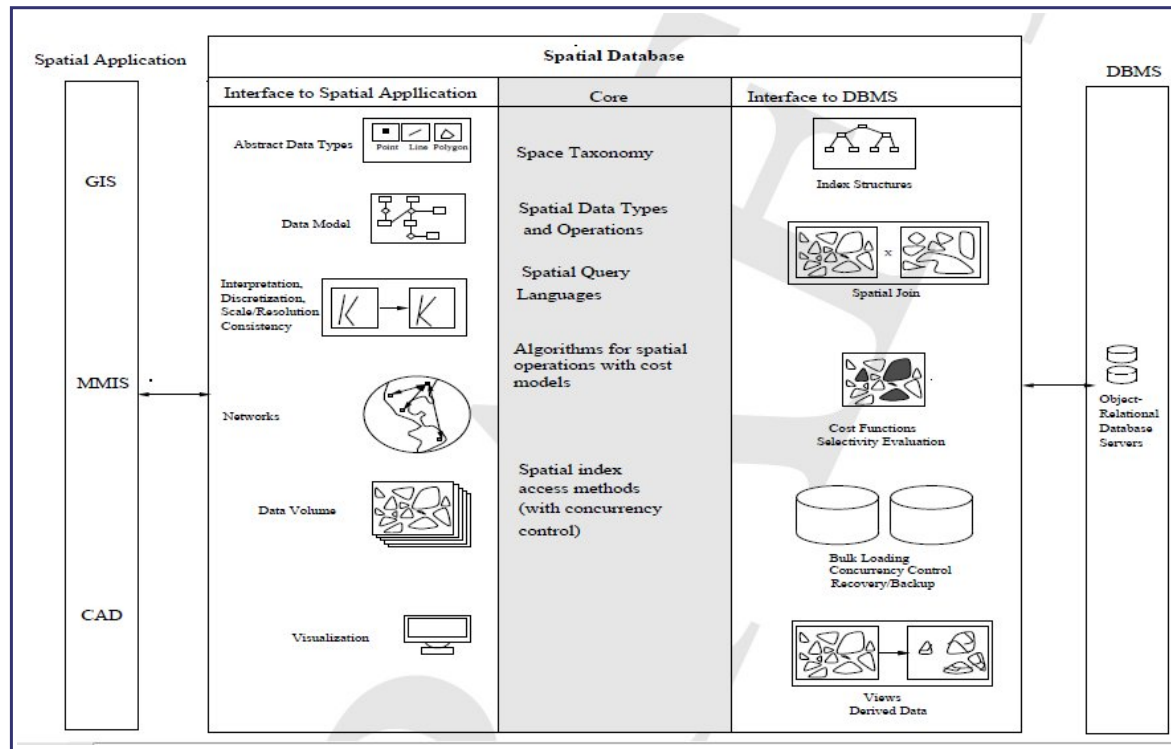




Look at this map, overlaid with scary data..

# SDBMS architecture





# GIS vs SDBMS

GIS is a specific application architecture built on top of a [more general purpose] SDBMS.


GIS typically tend to be used for:

<b>Search</b>	Thematic search, search by region, (re-)classification
<b>Location analysis</b>	Buffer, corridor, overlay
<b>Terrain analysis</b>	Slope/aspect, catchment, drainage network
<b>Flow analysis</b>	Connectivity, shortest path
<b>Distribution</b>	Change detection, proximity, nearest neighbor
<b>Spatial analysis/Statistics</b>	Pattern, centrality, autocorrelation, indices of similarity, topology: hole description
<b>Measurements</b>	Distance, perimeter, shape, adjacency, direction

# Spatial relationships

In 1D (and higher), spatial relationships can be expressed using 'intersects', 'crosses', 'within', 'touches' (these are T/F predicates).

Here is a sampling of spatial relationships in 2D:



Another diagram showing the [binary] operations:

IEEE Transactions on Knowledge and Data Engineering 6 (1): 86-95, 1994.

12


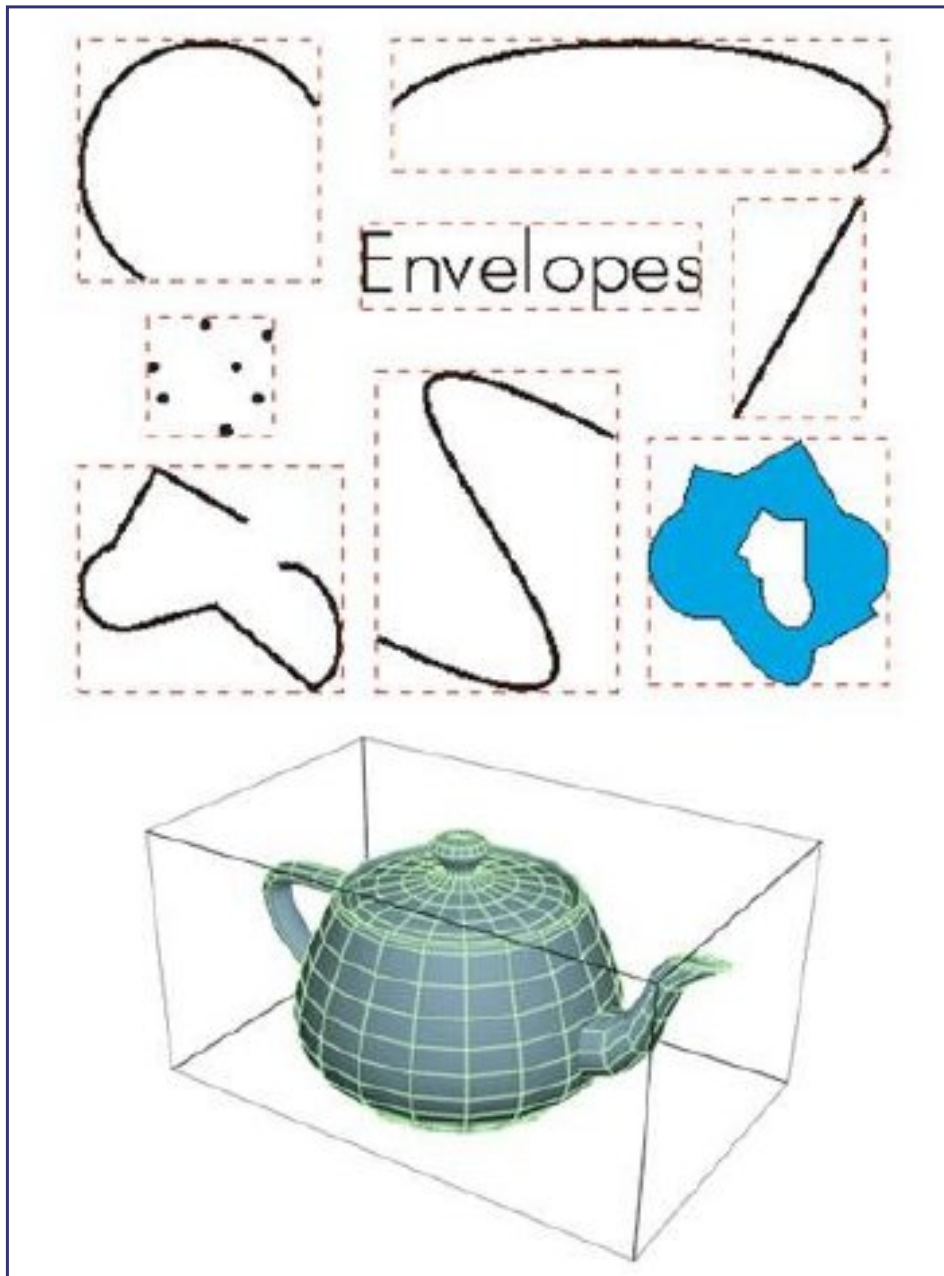


Figure 2. Examples of the binary topological relationships (a) disjoint, (b) meet, (c) overlap, (d) covers/covers\_by, (e) inside/contains, and (f) equal [24].

Minimum Bounding Rectangles (MBRs) are what are used to compute the results of operations shown above:



# Spatial relations – categories

Spatial relationships can be:

- topology-based [using defns of boundary, interior, exterior]
- metric-based [distance/Euclidian, angle measures]
- direction-based
- network-based [eg. shortest path]

Topological relationships could be further grouped like so:

- proximity
- overlap
- containment



# How can we put these relations to use?

We can perform the following, on spatial data:

- spatial measurements: find the distance between points, find polygon area..
- spatial functions: find nearest neighbors..
- spatial predicates: test for proximity, containment..

## Spatial Data Entity Creation

- Form an entity to hold county names, states, populations, and geographies

```
CREATE TABLE County(  
    Name      varchar(30),  
    State     varchar(30),  
    Pop       Integer,  
    Shape     Polygon);
```

## Spatial Data Entity Creation (Cont.)

- Form an entity to hold river names, sources, lengths, and geographies

```
CREATE TABLE River(  
    Name      varchar(30),  
    Source     varchar(30),  
    Distance   Integer,  
    Shape      LineString);
```

## Example Spatial Query

- Find all the counties that border on Contra Costa county

```
SELECT      C1.Name
FROM        County C1, County C2
WHERE       Touch(C1.Shape, C2.Shape) = 1
            AND C2.Name = 'Contra Costa';
```

## Example Spatial Query (Cont.)

- Find all the counties through which the Merced river runs

```
SELECT      C.Name, R.Name
FROM        County C, River R
WHERE       Intersect(C.Shape, R.Shape) = 1
            AND R.Name = 'Merced';
```

# Spatial operators, functions


## Spatial Functions

- Returns a geometry

- Union
- Difference
- Intersect
- XOR
- Buffer
- CenterPoint
- ConvexHull


- Returns a number

- LENGTH
- AREA
- Distance



# Spatial Operators

- Full range of spatial operators
  - Implemented as functional extensions in SQL
  - Topological Operators
    - Inside Contains
    - Touch Disjoint
    - Covers Covered By
    - Equal Overlap Boundary
  - Distance Operators
    - Within Distance
    - Nearest Neighbor



```
#query

+ equals(another :Geometry) : Boolean
+ disjoint(another :Geometry) : Boolean
+ intersects(another :Geometry) : Boolean
+ touches(another :Geometry) : Boolean
+ crosses(another :Geometry) : Boolean
+ within(another :Geometry) : Boolean
+ contains(another :Geometry) : Boolean
...

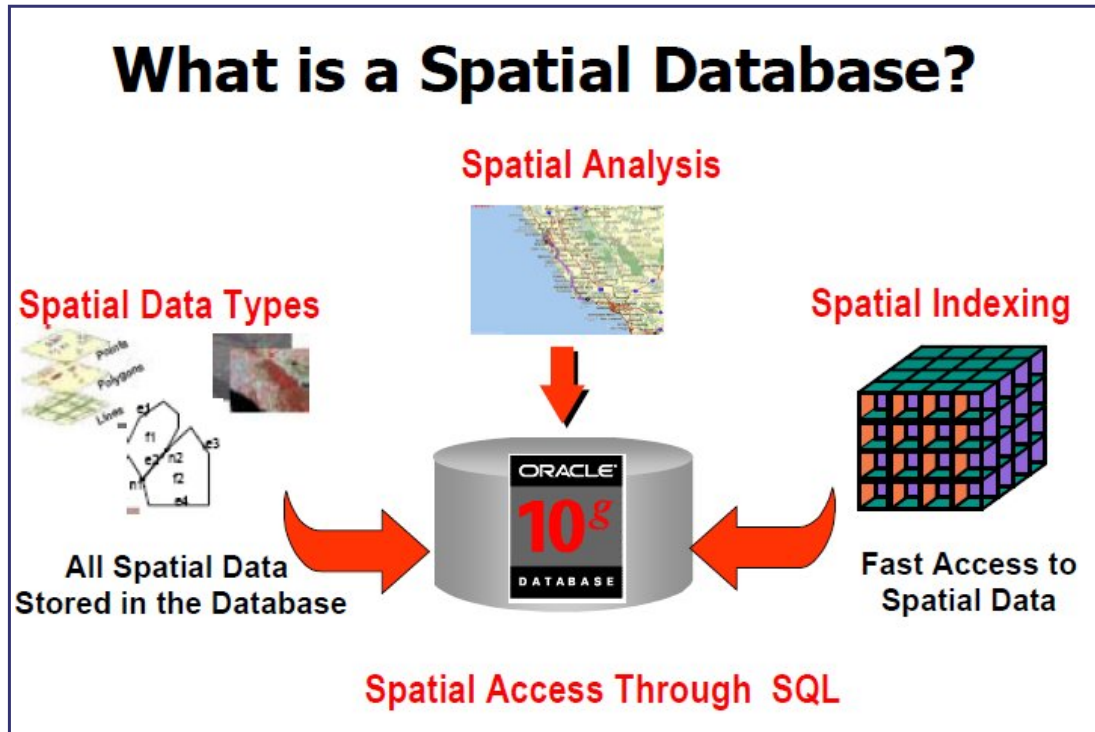
#analysis

+ distance(another : Geometry) : Distance
+ buffer(another : Distance) : Geometry
+ convexHull() : Geometry
...
```

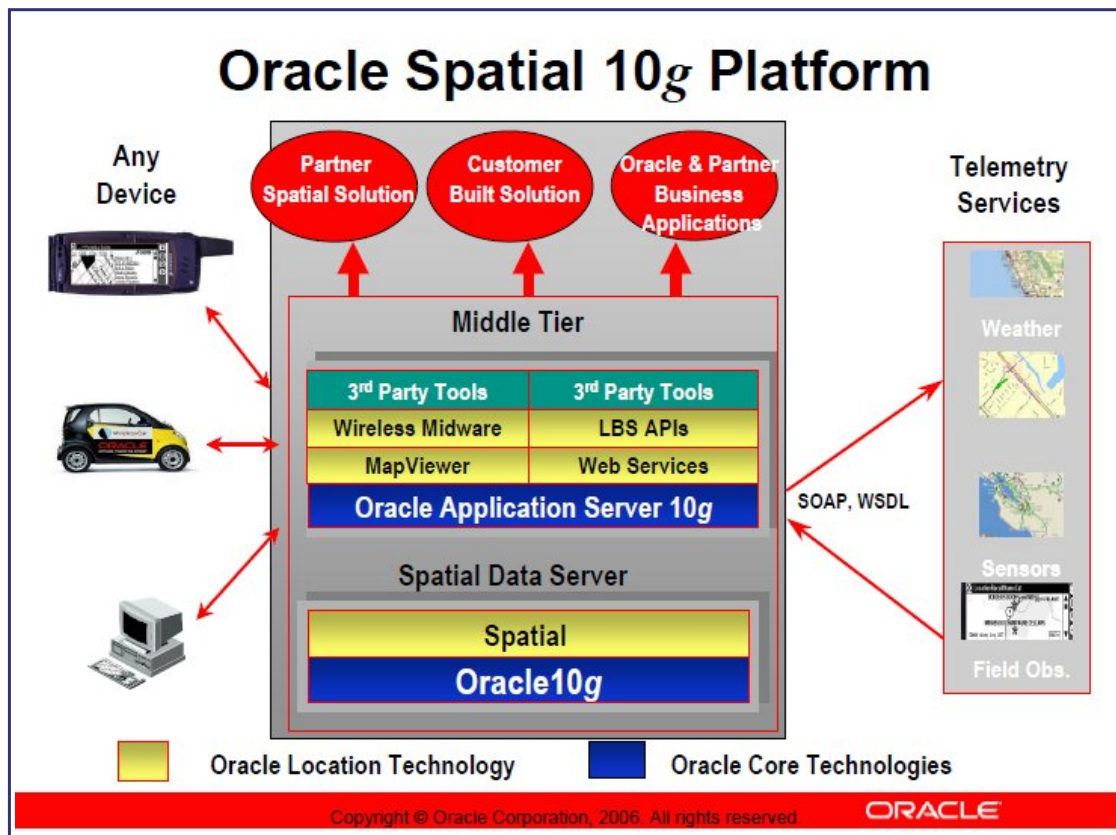
This doc [from 'FME Knowledge Center'; thanks to Minaxi Singla for the link] provides more info on the spatial operators.

# Oracle Spatial

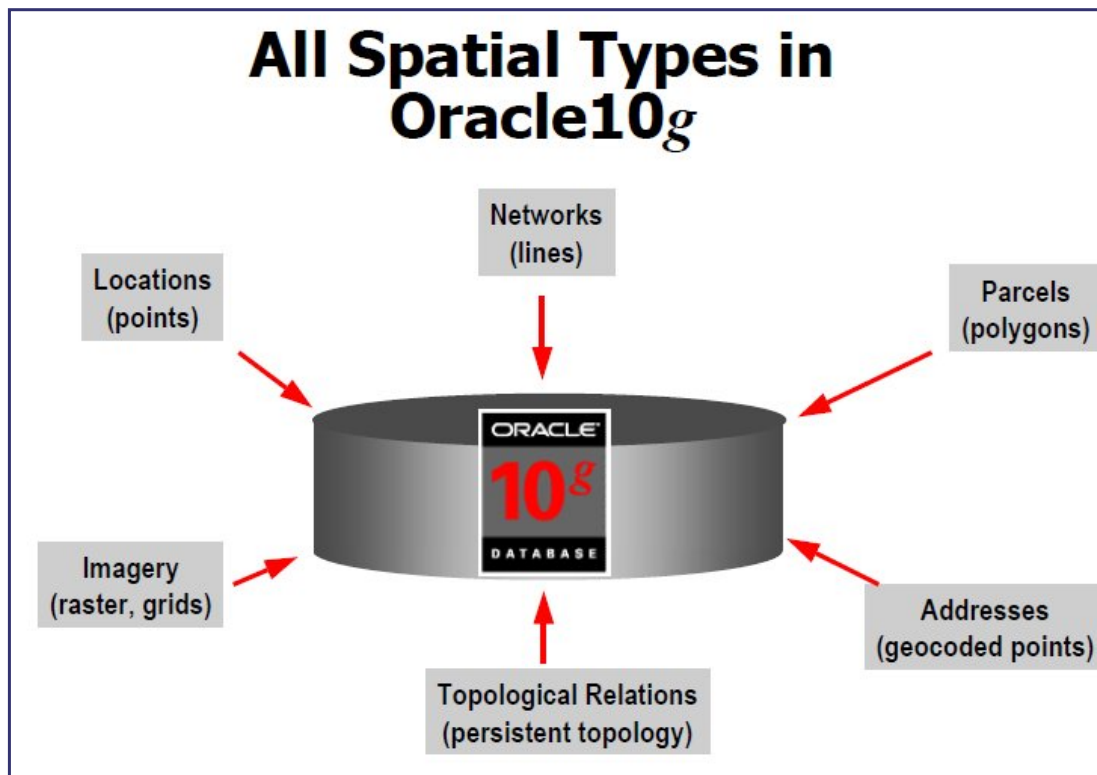
Oracle offers a 'Spatial' library for spatial queries – this includes UDTs and custom functions to process them.











# SDO\_GEOMETRY Object

- **SDO\_GEOMETRY** Object

SDO_GTYPE	NUMBER
SDO_SRID	NUMBER
SDO_POINT	SDO_POINT_TYPE
SDO_ELEM_INFO	SDO_ELEM_INFO_ARRAY
SDO_ORDINATES	SDO_ORDINATE_ARRAY

- Example

```
SQL> CREATE TABLE states (  
2     state      VARCHAR2(30) ,  
3     totpop     NUMBER(9) ,  
4     geom       SDO_GEOMETRY) ;
```

## SDO\_GEOMETRY Object

- **SDO\_GTYPE** – Defines the type of geometry stored in the object


GTYPE	Explanation
1 POINT	Geometry contains one point
2 LINESTRING	Geometry contains one line string
3 POLYGON	Geometry contains one polygon
4 HETEROGENEOUS COLLECTION	Geometry is a collection of elements of different types: points, lines, polygons
5 MULTIPOINT	Geometry has multiple points
6 MULTILINESTRING	Geometry has multiple line strings
7 MULTIPOLYGON	Geometry has multiple polygons

## SDO\_GTYPE

SDO_GTYPE	Four digit GTYPEs - Include dimensionality		
	2D	3D	4D
1 POINT	2001	3001	4001
2 LINESTRING	2002	3002	4002
3 POLYGON	2003	3003	4003
4 COLLECTION	2004	3004	4004
5 MULTIPOINT	2005	3005	4005
6 MULTILINESTRING	2006	3006	4006
7 MULTIPOLYGON	2007	3007	4007

# Constructing Geometries

```
SQL> INSERT INTO LINES VALUES (  
2>   attribute_1, ..., attribute_n,  
3>   SDO_GEOMETRY (  
4>     2002, null, null,  
5>     SDO_ELEM_INFO_ARRAY (1,2,1),  
6>     SDO_ORDINATE_ARRAY (  
7>       10,10, 20,25, 30,10, 40,10))  
8> );
```



# Spatial Operators

- Operators
  - **SDO\_FILTER**
    - Performs a primary filter only
  - **SDO\_RELATE** and **SDO\_<relationship>**
    - Performs a primary and secondary filter
  - **SDO\_WITHIN\_DISTANCE**
    - Generates a buffer around a geometry and performs a primary and optionally a secondary filter
  - **SDO\_NN**
    - Returns nearest neighbors

## SDO\_FILTER Example

- Find all the cities in a selected rectangular area
- Result is approximate

```
SELECT c.city, c.pop90
FROM proj_cities c
WHERE sdo_filter (
    c.location,
    sdo_geometry (2003, 32775, null,
        sdo_elem_info_array (1,1003,3),
        sdo_ordinate_array (1720300,1805461,
            1831559, 2207250))
) = 'TRUE';
```

Hint 1: All Spatial operators return TRUE or FALSE. When writing spatial queries always test with = 'TRUE', never <> 'FALSE' or = 'true'.

## SDO\_RELATE Example

- Find all counties in the state of New Hampshire

```
SELECT c.county, c.state_abrv
FROM geod_counties c,
     geod_states s
WHERE s.state = 'New Hampshire'
AND sdo_relate (c.geom,
    s.geom,
    'mask=INSIDE+COVEREDBY')
    = 'TRUE';
```

**Note:** For optimal performance, don't forget to index  
GEOD\_STATES(state)

## Relationship Operators Example

- Find all the counties around Passaic county in New Jersey:

```
SELECT /*+ ordered */ a.county
FROM geod_counties b,
     geod_counties a
WHERE b.county = 'Passaic'
      AND b.state = 'New Jersey'
      AND SDO_TOUCH(a.geom,b.geom) = 'TRUE';
```

- Previously:

```
. . .
AND SDO_RELATE(a.geom,b.geom,
               'MASK=TOUCH') = 'TRUE';
```



## SDO\_NN Example

- Find the five cities nearest to Interstate I170, ordered by distance

```
SELECT /*+ ordered */
      c.city, c.state_abrv,
      sdo_nn_distance (1) distance_in_miles
FROM   geod_interstates i,
      geod_cities c
WHERE  i.highway = 'I170'
      AND sdo_nn(c.location, i.geom,
                  'sdo_num_res=5 unit=mile', 1) = 'TRUE'
ORDER BY distance_in_miles;
```

- Note: Make sure you have an index on GEOD\_INTERSTATES (HIGHWAY).

## SDO\_WITHIN\_DISTANCE Examples

- Find all cities within a distance from an interstate

```
SELECT /*+ ordered */ c.city
FROM geod_interstates i, geod_cities c
WHERE i.highway = 'I170'
      AND sdo_within_distance (
          c.location, i.geom,
          'distance=15 unit=mile') = 'TRUE';
```

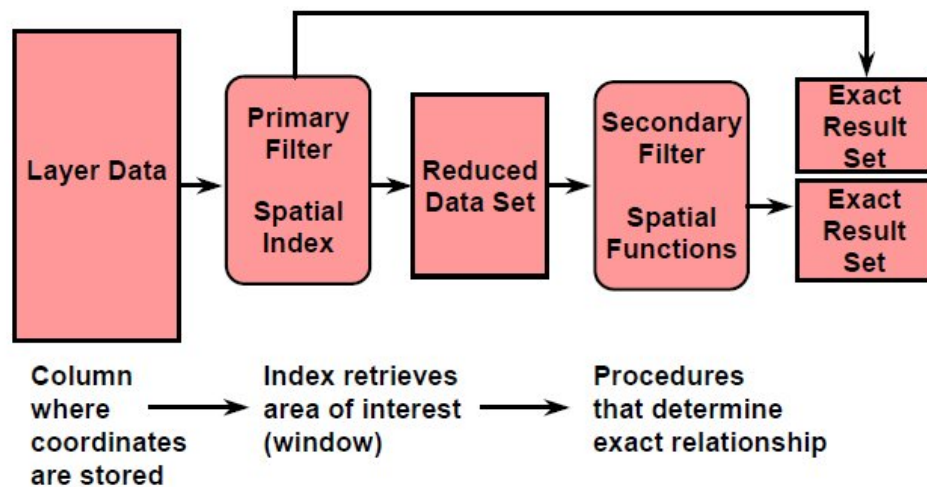
- Find interstates within a distance from a city

```
SELECT /*+ ordered */ i.highway
FROM geod_cities c, geod_interstates i
WHERE c.city = 'Tampa'
      AND sdo_within_distance (
          i.geom, c.location,
          'distance=15 unit=mile') = 'TRUE';
```

## Spatial Indexing

- Used to optimize spatial query performance
- R-tree Indexing
  - Based on minimum bounding rectangles (MBRs) for 2D data or minimum bounding volumes (MBVs) for 3D data
  - Indexes two, three, or four dimensions
- Provides an exclusive and exhaustive coverage of spatial objects
- Indexes all elements within a geometry including points, lines, and polygons

# Optimized Query Model



# Postgres PostGIS

## Types of queries - PostGIS

The function names for queries differ across geodatabases. The following list contains commonly used functions built into PostGIS, a free geodatabase which is a PostgreSQL extension (the term 'geometry' refers to a point, line, box or other two or three dimensional shape):

## **Types of queries - PostGIS (Cont.)**

1. Distance(geometry, geometry) : number
2. Equals(geometry, geometry) : boolean
3. Disjoint(geometry, geometry) : boolean
4. Intersects(geometry, geometry) : boolean
5. Touches(geometry, geometry) : boolean
6. Crosses(geometry, geometry) : boolean

## **Types of queries - PostGIS (Cont.)**

7. Overlaps(geometry, geometry) : boolean
8. Contains(geometry, geometry) : boolean
9. Intersects(geometry, geometry) : boolean
10. Length(geometry) : number
11. Area(geometry) : number
12. Centroid(geometry) : geometry

Here is an example – table creation, and polygon insertion:

```

Create Table County (
name VARCHAR(30),
shape geometry);
CREATE TABLE

Insert into County values ('Lynn', ST_Polygon(ST_GeomFromText('LINESTRING(75.15 29.53 1,77 29 1,77.6 29.5 1,
75.15 29.53 1)'),4326));
INSERT 0 1

SELECT *
FROM County;
name |
      shape
-----+-----
Lynn | 01030000A0E610000001000000040000009A99999999C9524048E17A14AE873D40000000000000F03F000000000040534000
00000000003D40000000000000F03F666666666665340000000000803D40000000000000F03F9A99999999C9524048E17A14AE873D
40000000000000F03F
(1 row)

Saty@Satys_USC_PC ~
$

```

To do the above, here are the steps on a PC (similar steps on a Mac):

- install Postgres (v.9.5, not 9.6 beta!)
- bring up 'Application Stack Builder' (an add-on that gets installed when Postgres v9.5 is installed), from the available installation options that come up, pick Spatial Extensions → 'PostGIS 2.2 for Postgres 9.5', install
- bring up a shell (I use 'cygwin'); note - if you want to use cygwin, be sure to use the shell that comes up when you run cygwin.bat, \*not\* the 'mintty' shell that you get when you double-click on the cygwin icon; Mac users would use the built-in shell
- 9.5/bin/initdb (on a Mac the path would be different)
- 9.5/bin/pg\_ctl start - this starts the Postgres server
- 9.5/bin/createdb mydb - a new db for us to create tables in
- 9.5/bin/psql.exe -d mydb -c "CREATE EXTENSION postgis;" - this adds spatial types to our db; note: 'psql' is the program that lets us communicate with the db server, via the shell
- 9.5/bin/psql.exe -d mydb -a -f county.sql - this is how you can execute SQL commands that you store in a .sql file

- edit the .sql file (eg add more data [including spatial data], create new tables, write SQL queries [including spatial ones]..), run the file (as shown above), edit, run.....
- 9.5/bin/pg\_ctl stop – optionally you can stop the server and restart it later
- ...

You can learn a lot about spatial queries from this page.

# Creating spatial indexes

As (more so than) with non-spatial data, the creation and use of spatial indexes VASTLY speed up processing!



# Can B Trees index spatial data?

In short, YES, if we pair it up with a 'z curve' indexing scheme (using a space-filling curve):

## Organizing spatial data with space filling curves

- Issue:

- Sorting is not naturally defined on spatial data
- Many efficient search methods are based on sorting datasets

- Space filling curves

- Impose an ordering on the locations in a multi-dimensional space
  - Examples: row-order (Fig. 1.11(a), z-order (Fig 1.11(b))
- Allow use of traditional efficient search methods on spatial data

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

(a)


7	8	14	16
5	6	13	15
2	4	10	12
1	3	9	11

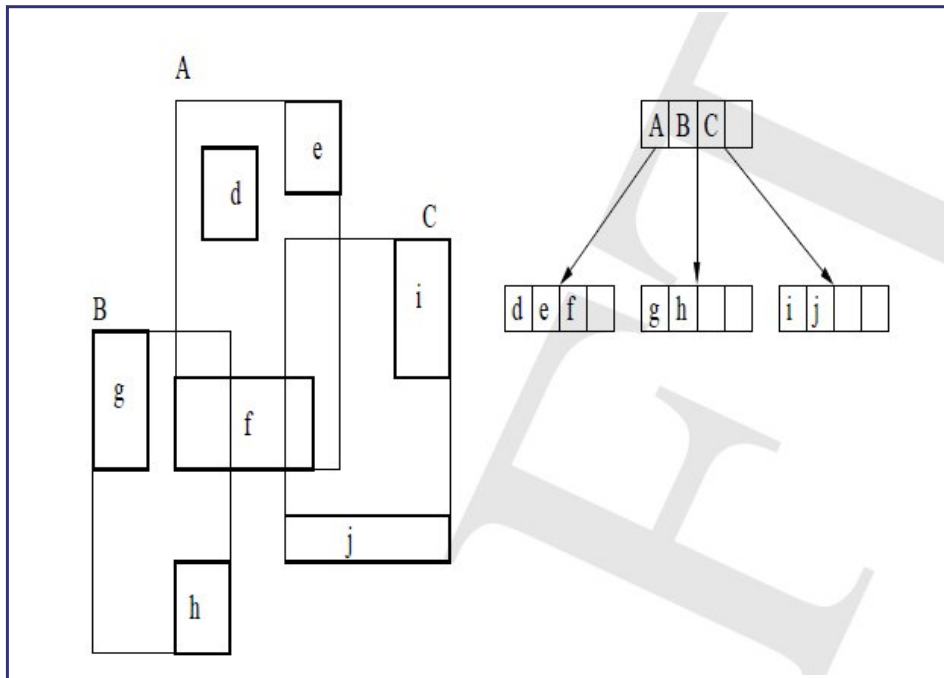
(b)

The idea is to quantize every (x,y) location into a recursively-divided 'quadtree' cell, and use the cell's binary (x,y) location to create a (binary) 'z' key, which is ordered along the unit (0..1) interval – in other words, 2D (x,y) points get mapped (indexed) to ordered 1D 'z' locations.

But, this is of academic interest mostly, not commonly practiced in industry – Apple's FoundationDB is an exception.

# R trees






R trees use MBRs to create a hierarchy of bounds.


Variations, FYI: R+ tree, R\* tree, Buddy trees, Packed R trees..

# k-d trees, K-D-B trees


k-d tree



Alternate: K-D-B tree:




(a) Planar partition




(b) A hierarchical K-D-B-tree structure

# Quadrees (and octrees)

Here we recursively and adaptively subdivide space [subdivisions happen only where necessary].






A representation of how a quadtree divides an indexed area. Source: [Wikipedia](#)

Each node is either a leaf node, with indexed points or null, or an internal (non-leaf) node that has exactly 4 children. The hierarchy of such nodes forms the quadtree.



# Indexing evolution



Indexing schemes continue to evolve.

# Query processing: filter, refine

## Query Processing

- Efficient algorithms to answer spatial queries
- Common Strategy - filter and refine
  - Filter Step: Query Region overlaps with MBRs of B, C and D
  - Refine Step: Query Region overlaps with B and C


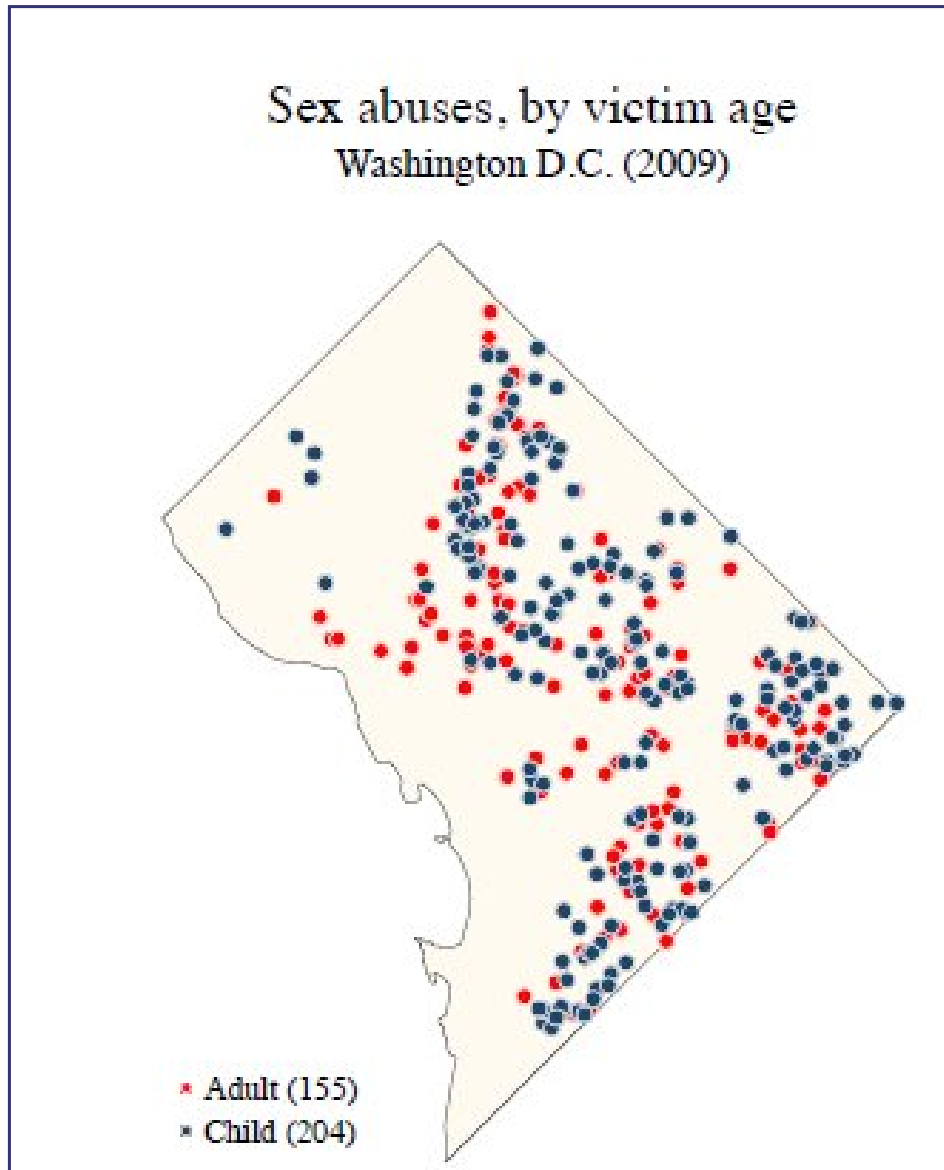


Fig 1.8

# Visualizing spatial data

A variety of non-spatial attrs can be mapped on to spatial data, providing an intuitive grasp of patterns, trends and abnormalities. Following are some examples.

Dot map:



Here's another one.

Proportional symbol map:

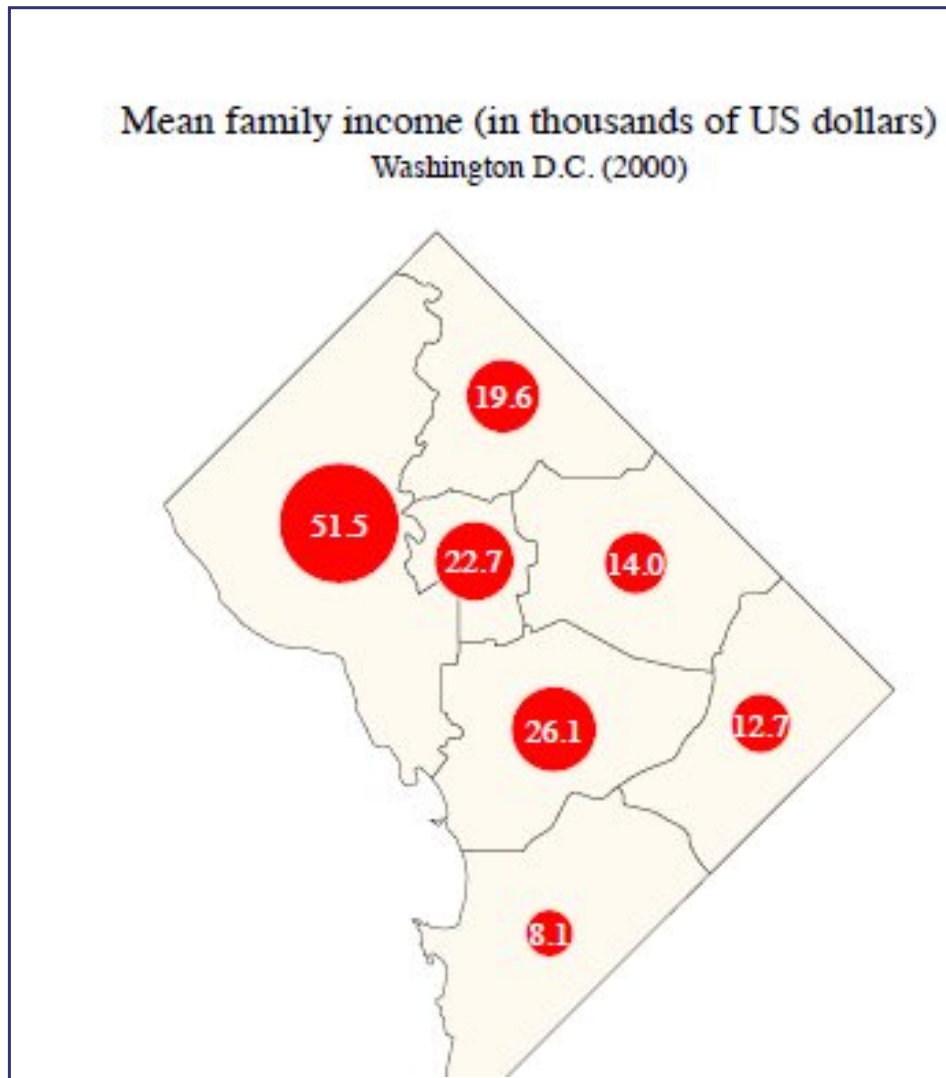
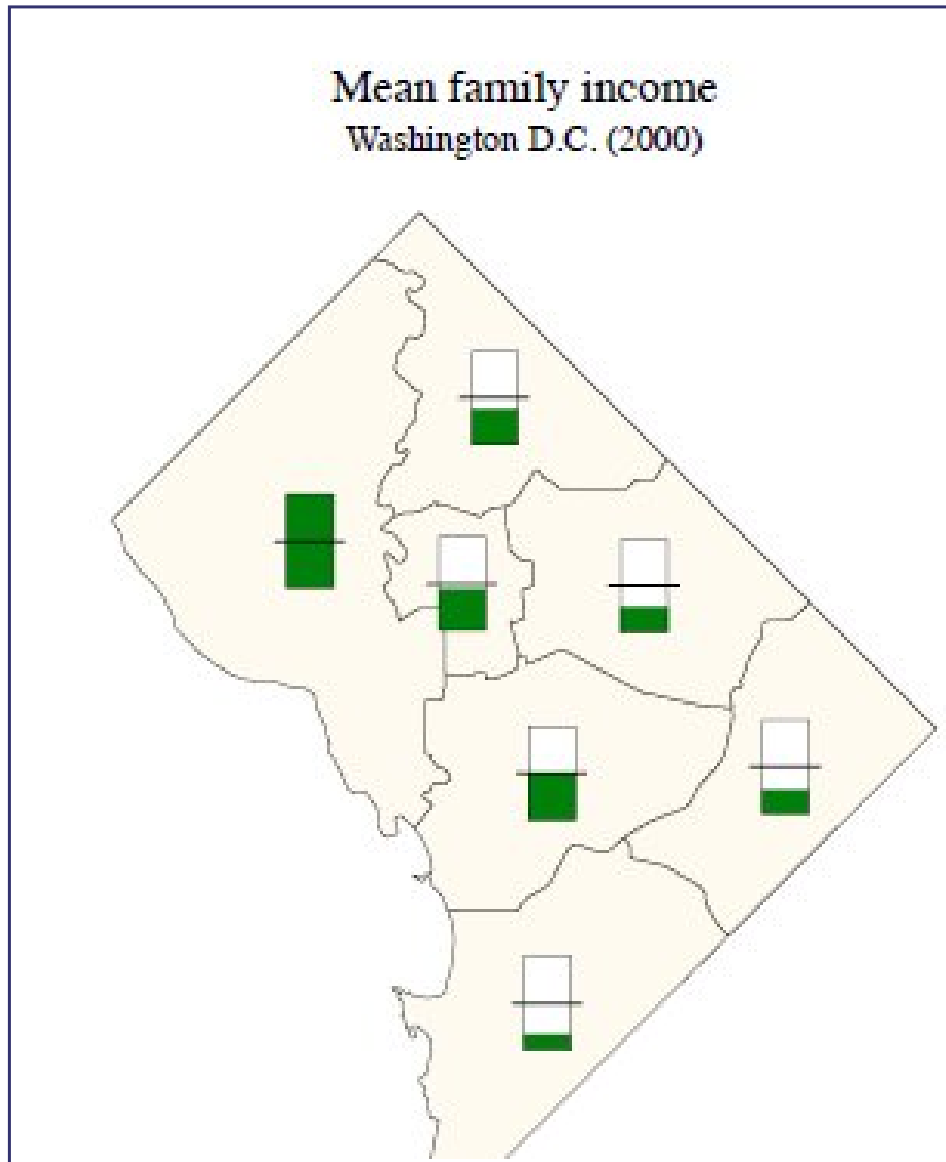
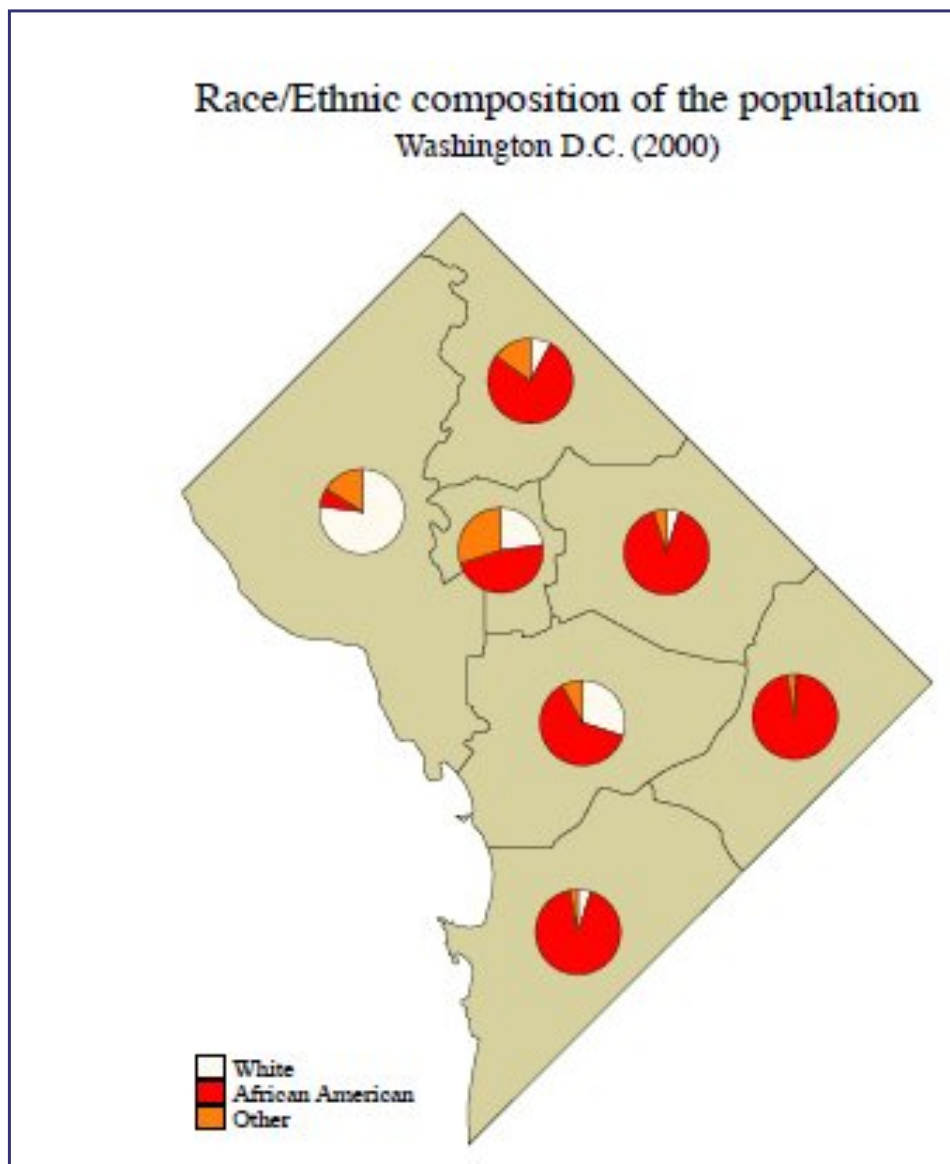


Diagram map:

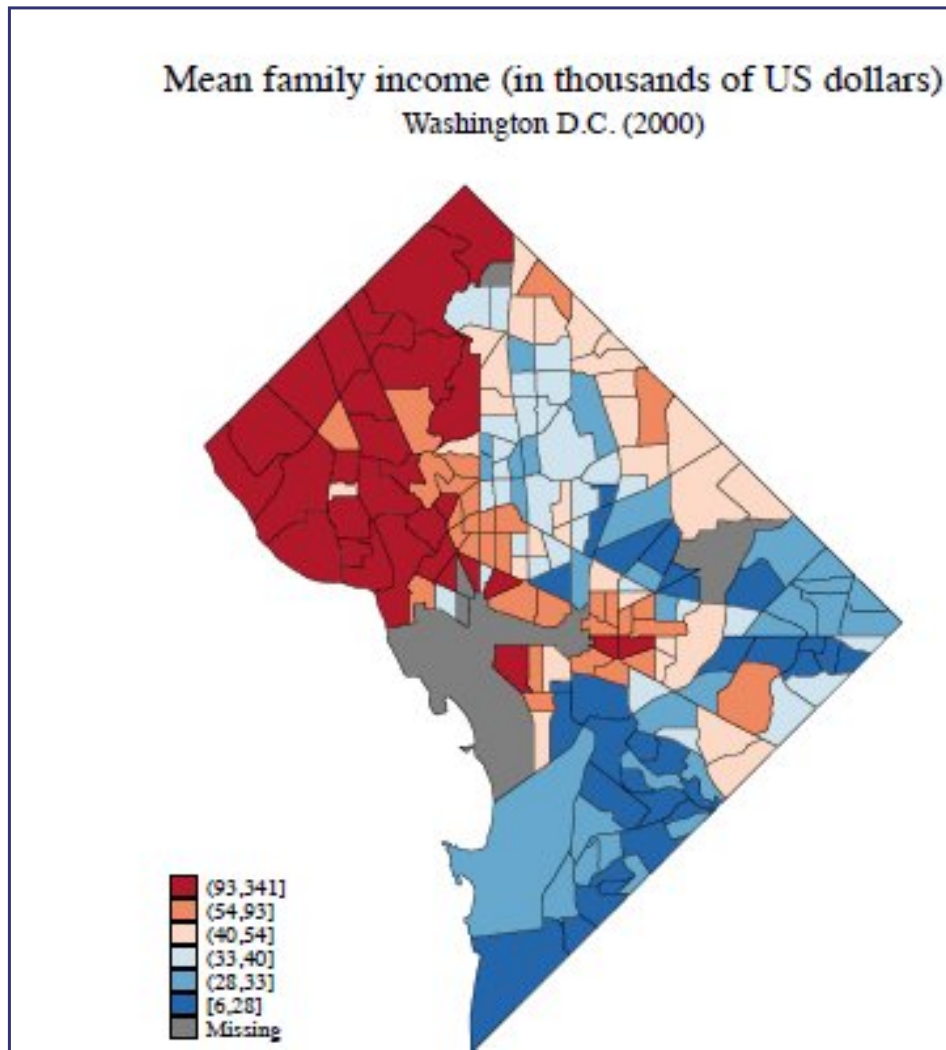


Another diagram map:

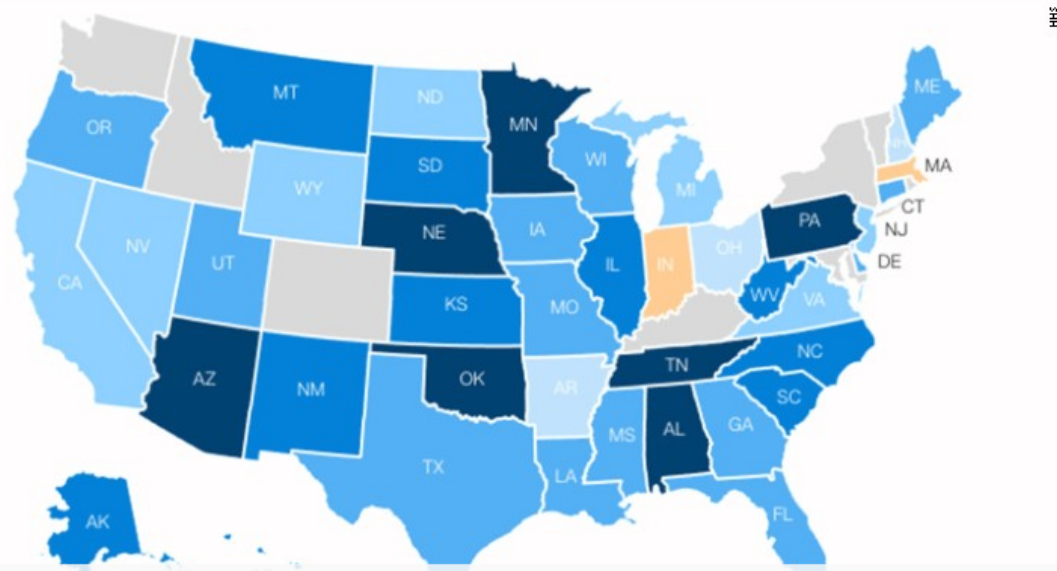


Also possible to plot multivariate data this way.

Choropleth maps (plotting of a variable of interest, to cover an entire region of a map):



# Obamacare's true cost



Most won't experience the full price hike or anything near it



# So who (else) has spatial extensions?

Everyone!

Thanks to SQL's facility for custom datatype ('UDT') and function creation ('functional extension'), "spatial" has been implemented for every major DB out there:

- Oracle: Locator, Spatial, SDO
- Postgres: PostGIS
- DB2: Spatial Datablade
- Informix: Geodetic Datablade
- SQL Server: Geometric and Geodetic Geography types
- MySQL: spatial library comes 'built in'
- SQLite: SpatiaLite
- ..

# Google KML

Google's KML format is used to encode spatial data for Google Earth, etc. [Here is a page on importing other geospatial dataset formats into Google Earth.](#)

# OpenLayers

OpenLayers is an open GIS platform.

# ESRI: Arc\*

.....  
ESRI is the home of the powerful, flexible family of ArcGIS products – and they are local!

# QGIS etc.

There is a variety of inexpensive/open source mapping platforms, competing with more pricey commercial offerings (from ESRI etc). Here are several:

- QGIS
- MapBox
- Carto
- GIS Cloud

